

Parallel Programming HW4

Blocked All-Pairs Shortest Path

Implementation

- Single GPU


1. Change `Dist[]` to 1-D array
2. Use `dim3` data structure to launch GPU
3. Inside kernel function, use `blockIdx.x` and `blockIdx.y` to relabel block indices, and `threadIdx.x` and `threadIdx.y` to relabel real array indices, so such four loops can be parallelized
4. After each `k`, threads should be synchronized so the afterward iterations have the updated values

- OpenMp


1. The mentioned modifications from single GPU version
2. Change `dev_dist[]` to a 2-D array that GPU0 holds `dev_dist[0]` and GPU1 holds `dev_dist[1]`
3. Inside `block_APSP()`, use `omp_get_thread_num()` as `gpu_id` to realize the parallelism
4. Phase 1: Both GPU calculates the pivot block
Phase 2: GPU0 gets the pivot row, and the upper part of pivot column
GPU1 gets the pivot row, and the lower part of the pivot column
Phase 3: GPU0 gets the upper part of the `dev_dist[]`
GPU1 gets the lower part of the `dev_dist[]`
5. After phase 3, GPU0 and GPU1 have to copy the upper part and the lower part respectively to host, and take the complementary part back to GPU.

- MPI


1. All phases are same as OpenMp version
2. After each round, GPU0 and GPU1 have to copy the upper part and the lower part respectively to host
3. Process0 sends the upper part to process1, and process1 sends the lower part to process0 using non-blocking send
4. Each process copies the received part to each GPU



For OpenMp and MPI versions, GPU0 and GPU1 get the pivot block and do calculation in phase 1.



In phase 2, GPU0 gets the pivot row and the upper pivot column; GPU1 also gets the pivot row, and the lower pivot column.



In phase 3, GPU0 gets the upper part above the pivot row; GPU1 gets the lower part below the pivot row.

Profiling Results

- nvprof ./HW4_cuda.exe Testcase/in5 out5

```
[user12@gpucluster2 HW4]$ nvprof ./HW4_cuda.exe Testcase/in5 out5
==26999== NVPROF is profiling process 26999, command: ./HW4_cuda.exe Testcase/in5 out5
total time 59101.933
==26999== Profiling application: ./HW4_cuda.exe Testcase/in5 out5
==26999== Profiling result:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
99.05%	56.3268s	846	66.580ms	934ns	569.32ms	cal(int, int, int, int, int, int, int, int, int, int*)
0.50%	282.34ms	1	282.34ms	282.34ms	282.34ms	[CUDA memcpy DtoH]
0.45%	257.03ms	1	257.03ms	257.03ms	257.03ms	[CUDA memcpy HtoD]

This part is the main results I use as the experiment outcomes

```
==26999== API calls:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
99.43%	58.6655s	2	29.3327s	257.67ms	58.4078s	cudaMemcpy
0.52%	307.08ms	1	307.08ms	307.08ms	307.08ms	cudaMalloc
0.03%	20.630ms	1	20.630ms	20.630ms	20.630ms	cudaFree
0.01%	5.5659ms	846	6.5790us	6.0720us	39.422us	cudaLaunch
0.00%	1.8420ms	6768	272ns	246ns	12.048us	cudaSetupArgument
0.00%	510.44us	166	3.0740us	261ns	108.15us	cuDeviceGetAttribute
0.00%	312.43us	846	369ns	338ns	3.0700us	cudaConfigureCall
0.00%	79.953us	2	39.976us	35.628us	44.325us	cuDeviceTotalMem
0.00%	58.336us	2	29.168us	26.152us	32.184us	cuDeviceGetName
0.00%	4.0650us	1	4.0650us	4.0650us	4.0650us	cudaSetDevice
0.00%	2.1490us	2	1.0740us	539ns	1.6100us	cuDeviceGetCount
0.00%	1.8490us	4	462ns	337ns	571ns	cuDeviceGet
0.00%	1.4800us	1	1.4800us	1.4800us	1.4800us	cudaGetDeviceCount

- nvprof ./HW4_openmp.exe Testcase/in5 out5

```
[user12@gpucluster2 HW4]$ nvprof ./HW4_openmp.exe Testcase/in5 out5
==27241== NVPROF is profiling process 27241, command: ./HW4_openmp.exe Testcase/in5 out5
total time 66265.210
==27241== Profiling application: ./HW4_openmp.exe Testcase/in5 out5
==27241== Profiling result:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
81.53%	57.8522s	1128	51.287ms	793ns	573.53ms	cal(int, int, int, int, int, int, int, int, int, int*)
9.70%	6.87938s	189	36.399ms	195.62us	703.57ms	[CUDA memcpy HtoD]
8.77%	6.22440s	188	33.109ms	178.24us	632.93ms	[CUDA memcpy DtoH]

1. computing time
2. memcpy time
3. memcpy time

```
==27241== API calls:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
88.75%	70.1226s	188	372.99ms	1.0760us	1.25781s	cudaMemcpyPeer
4.11%	3.25028s	3	1.08343s	650.47ms	1.89358s	cudaMemcpy
3.72%	2.93531s	2	1.46766s	1.08206s	1.85325s	cudaMalloc
3.42%	2.69864s	1128	2.3924ms	6.5090us	496.89ms	cudaLaunch
0.00%	3.2953ms	9024	365ns	246ns	12.572us	cudaSetupArgument
0.00%	669.12us	1128	593ns	386ns	13.410us	cudaConfigureCall
0.00%	633.45us	191	3.3160us	1.2660us	19.336us	cudaSetDevice
0.00%	578.00us	166	3.4810us	264ns	141.42us	cuDeviceGetAttribute
0.00%	88.225us	2	44.112us	36.663us	51.562us	cuDeviceGetName
0.00%	76.671us	2	38.335us	35.557us	41.114us	cuDeviceTotalMem
0.00%	37.867us	1	37.867us	37.867us	37.867us	cudaFree
0.00%	3.1590us	4	789ns	413ns	929ns	cuDeviceGet
0.00%	2.2240us	2	1.1120us	670ns	1.5540us	cuDeviceGetCount

- `nvprof --print-summary-per-gpu ./HW4_omp.exe Testcase/in5 out5`

```
[user12@gpucluster2 HW4]$ nvprof --print-summary-per-gpu ./HW4_omp.exe Testcase/in5 out5
==27495== NVPROF is profiling process 27495, command: ./HW4_omp.exe Testcase/in5 out5
total    time 56900.494
==27495== Profiling application: ./HW4_omp.exe Testcase/in5 out5
==27495== Profiling result:

==27495== Device "Tesla M2090 (0)"
Time(%)   Time      Calls      Avg      Min      Max  Name
92.21%   28.9818s      564   51.386ms   845ns   575.07ms  cal(int, int, int, int, int, int, int, int, int*)
 3.91%    1.22786s       94   13.062ms  239.46us   69.300ms  [CUDA memcpy DtoH]
 3.88%    1.21986s       95   12.841ms  194.08us   61.025ms  [CUDA memcpy HtoD]

==27495== Device "Tesla M2090 (1)"
Time(%)   Time      Calls      Avg      Min      Max  Name
92.32%   28.8798s      564   51.205ms   782ns   569.17ms  cal(int, int, int, int, int, int, int, int, int*)
 3.98%    1.24388s       94   13.233ms  261.38us   83.971ms  [CUDA memcpy HtoD]
 3.71%    1.15996s       94   12.340ms  177.89us   24.607ms  [CUDA memcpy DtoH]
```

Information about different threads can also be separately output

- `mpirun -np 2 -hostfile hostfile ./profile.sh -o result.%q[PMI_RANK] ./HW4_mpi.exe Testcase/in5 out5`

`nvprof --import-profile result.0.`

```
[user12@gpucluster2 HW4]$ mpirun -np 2 -hostfile hostfile ./profile.sh -o result.%q[PMI_RANK] ./HW4_mpi.exe Testcase/in5 out5
==27767== NVPROF is profiling process 27767, command: ./HW4_mpi.exe Testcase/in5 out5
==27765== NVPROF is profiling process 27765, command: ./HW4_mpi.exe Testcase/in5 out5
total[0]  time 101235.293
comp[0]   time 30521.711
comm[0]   time 48565.561
mem[0]    time 21678.266
total[1]  time 101284.605
comp[1]   time 29922.008
comm[1]   time 48075.114
mem[1]    time 22896.766
==27765== Generated result file: /home/cs542200/user12/HW4/result.0.
==27767== Generated result file: /home/cs542200/user12/HW4/result.1.
[user12@gpucluster2 HW4]$ nvprof --import-profile result.0.
===== Profiling result:
Time(%)   Time      Calls      Avg      Min      Max  Name
58.21%   29.0014s      564   51.421ms   849ns   568.17ms  cal(int, int, int, int, int, int, int, int, int*)
27.01%   13.4587s       95   141.67ms  441.99us   2.49106s  [CUDA memcpy HtoD]
14.78%    7.36277s       93   79.170ms  251.88us   490.58ms  [CUDA memcpy DtoH]

===== API calls:
Time(%)   Time      Calls      Avg      Min      Max  Name
58.36%   30.6486s       94   326.05ms  16.682ms   616.00ms  cudaEventSynchronize
41.38%   21.7348s      189   115.00ms    964ns   2.52154s  cudaMemcpy
 0.22%    1.1750ms        1   115.25ms  115.25ms   115.25ms  cudaMalloc
 0.02%    0.371ms       166    62.477us   354ns    8.8386ms  cuDeviceGetAttribute
 0.01%    6.8556ms      564   12.155us   6.4140us   57.975us  cudaLaunch
 0.00%    1.7430ms     4512    386ns    246ns    11.604us  cudaSetupArgument
 0.00%    1.1750ms      188    6.2500us   3.0830us   19.368us  cudaEventRecord
 0.00%    418.32us       94    4.4500us   2.3800us   15.869us  cudaEventElapsedTime
 0.00%    379.11us      564    672ns    360ns    9.1090us  cudaConfigureCall
 0.00%    290.50us        1   290.50us   290.50us   290.50us  cudaFree
 0.00%    257.14us        2   128.57us   100.82us   156.32us  cuDeviceTotalMem
 0.00%    121.54us        2    60.768us   58.304us   63.232us  cuDeviceGetName
 0.00%    13.948us        2    6.9740us   1.9720us   11.976us  cudaEventCreate
 0.00%    6.4300us        1    6.4300us   6.4300us   6.4300us  cudaSetDevice
 0.00%    3.0000us         4    750ns     616ns     874ns    cuDeviceGet
 0.00%    2.9580us        2    1.4790us   956ns     2.0020us  cuDeviceGetCount
```

nvprof —import-profile result.1.

```
[user12@gpucluster2 HW4]$ nvprof --import-profile result.1.
```

===== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
56.45%	28.9063s	564	51.252ms	946ns	569.74ms	cal(int, int, int, int, int, int, int, int, int*)
22.88%	11.7191s	94	124.67ms	203.40us	1.60258s	[CUDA memcpy DtoH]
20.67%	10.5840s	94	112.60ms	257.92us	843.64ms	[CUDA memcpy HtoD]

===== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
56.41%	29.9467s	94	318.58ms	8.8569ms	946.22ms	cudaEventSynchronize
43.27%	22.9696s	189	121.53ms	1.3450us	1.60381s	cudaMemcpy
0.30%	160.40ms	1	160.40ms	160.40ms	160.40ms	cudaMalloc
0.01%	6.5765ms	564	11.660us	6.5140us	58.254us	cudaLaunch
0.00%	1.7231ms	4512	381ns	256ns	8.3400us	cudaSetupArgument
0.00%	1.6952ms	166	10.212us	330ns	501.49us	cuDeviceGetAttribute
0.00%	1.0909ms	188	5.8020us	2.9900us	13.097us	cudaEventRecord
0.00%	373.84us	94	3.9770us	2.5830us	10.947us	cudaEventElapsedTime
0.00%	362.13us	1	362.13us	362.13us	362.13us	cudaFree
0.00%	348.10us	564	617ns	327ns	6.8860us	cudaConfigureCall
0.00%	176.90us	2	88.448us	85.979us	90.917us	cuDeviceTotalMem
0.00%	110.90us	2	55.447us	45.886us	65.009us	cuDeviceGetName
0.00%	11.096us	2	5.5480us	1.9600us	9.1360us	cudaEventCreate
0.00%	6.4870us	1	6.4870us	6.4870us	6.4870us	cudaSetDevice
0.00%	3.3410us	2	1.6700us	1.0400us	2.3010us	cuDeviceGetCount
0.00%	2.9530us	4	738ns	630ns	837ns	cuDeviceGet


Experiment & Analysis

• Time Evaluation:

- Always select the maximum time as the termination of execution.

• GFLOPS Calculation:

- I count the total calculation required in one execution as the total amount of floating operations, then divide it by 10^9 and the execution time.
- For single GPU version, the counting is done as below:
 - (1) An execution requires $\frac{n}{B}$ rounds
 - (2) During each round, there are $\left(\frac{n}{B}\right)^2$ blocks
 - (3) To do APSP, there are B intermediate vertices
 - (4) For each intermediate vertex, B^2 calculations are required
 - (5) Obtain counts of calculation is $\frac{n}{B} \times \left(\frac{n}{B}\right)^2 \times B \times B^2 = n^3$
- For OpenMp and MPI versions, the counting is done as below:
 - (1) Consider GPU0, and recall that GPU0 gets the upper part of distance matrix
 - (2) For round i, the number of blocks is $i \times \frac{n}{B}$, so total number of blocks calculated after execution is $\frac{n}{B} \sum_{i=1}^{n/B} i = \frac{n^2 B + n^3}{2B^3}$
 - (3) Additionally, B intermediate vertices and B^2 calculations for each
 - (4) Obtain counts of calculation is $\frac{n^2 B + n^3}{2B^3} \times B \times B^2 = \frac{n^2(B+n)}{2}$



For each round in OpenMp and MPI versions, GPU0 gets the blue parts and GPU1 get the red parts, and it's evident that the number of blocks increases row by row.

• Bandwidth Calculation:

- Regard the number of accessing `dev_dist[]` per kernel runtime as the expected bandwidth.
- For simplicity, an approximation is done by increment the number of accessing `dev_dist[]` by 6 for the distance update section for each calculation.
- Obtain the bandwidth by multiplying GFLOPS by 6, and by 4 for the unit GB, so the plots of GFLOPS and bandwidth are basically the same.

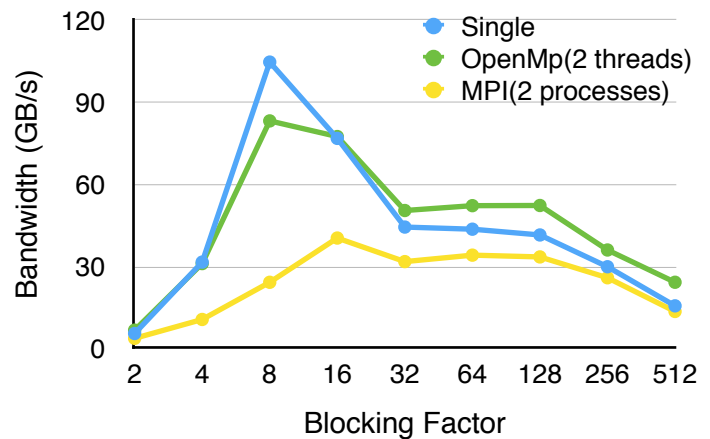
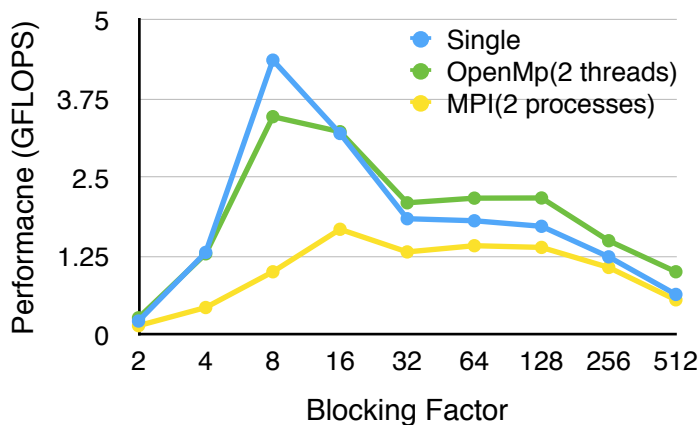
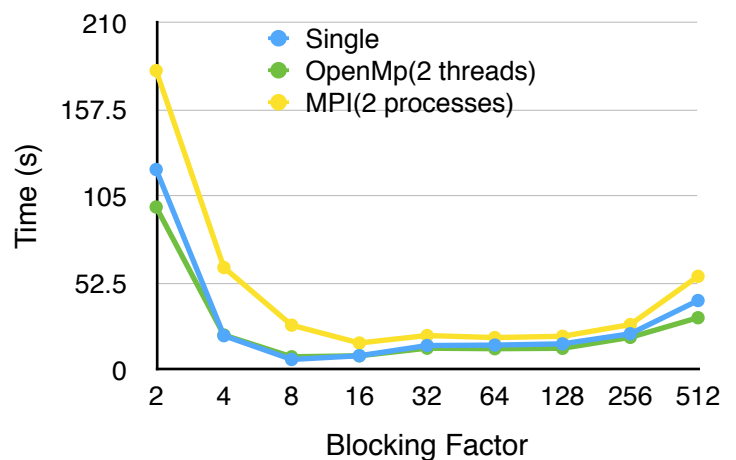
• Execution Time & GFLOPS & Bandwidth v.s. Blocking Factor I

- Input Setting :

```
block(round, round)
thread(min(32,B), min(32,B))
Testcase in4
OpenMp (2 threads)
MPI (2 processes)
```

- Results :

	Single	OpenMp (2 threads)	MPI (2 processes)
B=2	121.139329	98.429806	180.994125
B=4	20.567450	20.953182	61.823029
B=8	6.183443	7.804682	26.938044
B=16	8.432997	8.404206	16.154587
B=32	14.611323	12.990598	20.665523
B=64	14.872236	12.680870	19.430720
B=128	15.630576	12.929623	20.223118
B=256	21.735550	19.562106	27.286578
B=512	41.839919	31.440321	56.492201



- Observation :

- (1) For B=8~128, the execution time remains relatively low. When blocking factor is too small (B=2) or too large (B=512), the execution time tends to be longer.
- (2) For B<32, the GFLOPS is hard to reveal conclusion because the dimensions of block and thread vary from B to B. For B≥32, since the dimension of thread is fixed, the tendency is more reasonable.
- (3) To get a clearer view, I fixed the dimensions of block and thread in the next experiment.



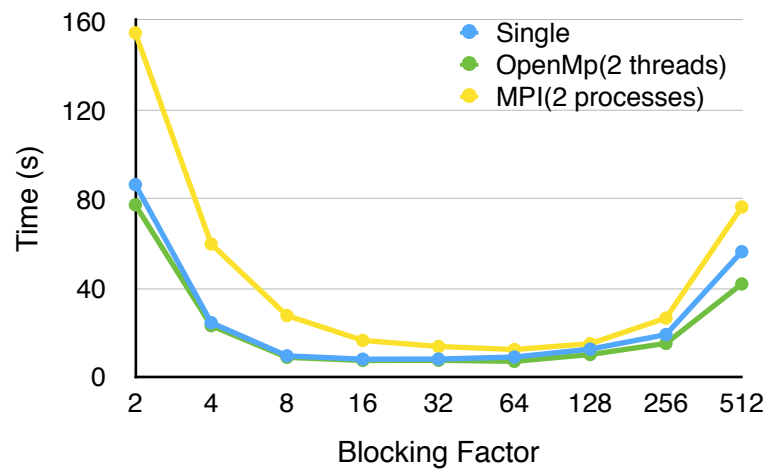
• Execution Time & GFLOPS & Bandwidth v.s. Blocking Factor II

- Input Setting :

block(10, 10)
 thread(10, 10)
 Testcase in4
 OpenMp (2 threads)
 MPI (2 processes)

- Results :

	Single	OpenMp (2 threads)	MPI (2 processes)
B=2	86.327791	77.210647	154.572869
B=4	24.26173	22.904621	59.636753
B=8	9.336398	8.663839	27.456777
B=16	7.835076	7.248722	16.272898
B=32	7.897341	7.346052	13.551727
B=64	8.787723	6.864621	12.190754
B=128	12.353731	9.932101	14.808557
B=256	18.883344	14.911443	26.315696
B=512	56.171169	41.661553	76.245637



- Observation :

- (1) MPI version requires longer execution time because it suffers from communication between processes. Thus the GFLOPS of MPI version is the lowest among the three.


- (2) OpenMp version is the fastest and achieves higher performance than the others. It is only slightly faster than single version because after each round, GPUs have to exchange information by `cudaMemcpy()`.
- (3) When B is too small, for example, B=2, the blocking is not effective because in a 3000x3000 matrix, a 2x2 block is almost as small as the non-blocked size. It certainly helps, but not much though. Therefore, when B=2 and 4, a large amount of time is still needed.
- (4) As blocking factor grows, more elements are being parallelized, so the execution time decreases.
- (5) When B is too large, the innate limitation of hardware may cause slowdown. During phase 3, three blocks are to be accessed, but if the block size is too large that all three of such blocks do not fit into the GPU cache, then cache misses occur, which considerably influences the performance.
(Gayathri Venkataraman, Sartaj Sahni, and Srabani Mukhopadhyaya, "A Blocked All-Pairs Shortest-Paths Algorithm")

- Time Distribution

- Input Setting :

```
block(round, round)
thread(min(32,B), min(32,B))
B = 64
OpenMp (2 threads)
MPI (2 processes)
```

- Results :



- Observation :

- (1) The computation time of OpenMp and MPI versions is half of that of single version, and this is reasonable because the problem is distributed to 2 GPUs in OpenMp and MPI.
- (2) Although MPI can do computation in parallel, it suffers from massive amount of communication overhead that dominates the overall performance. Such overhead leads to the consequences in the previous experiments that MPI version requires longer execution time.
- (3) Computing time increases as the input size grows.
- (4) Communication time increases as the input size grows because there are more data to be sent and received.
- (5) Memory copy time also increases as input size grows because there are more data to be transferred between host and device. But OpenMp and MPI have the same amount of memory copy time because the times of memory copy calls are the same while single version calls `cudaMemcpy()` only twice.

- Weak Scalability & Optimization (Pinned Memory)

- Input Setting :

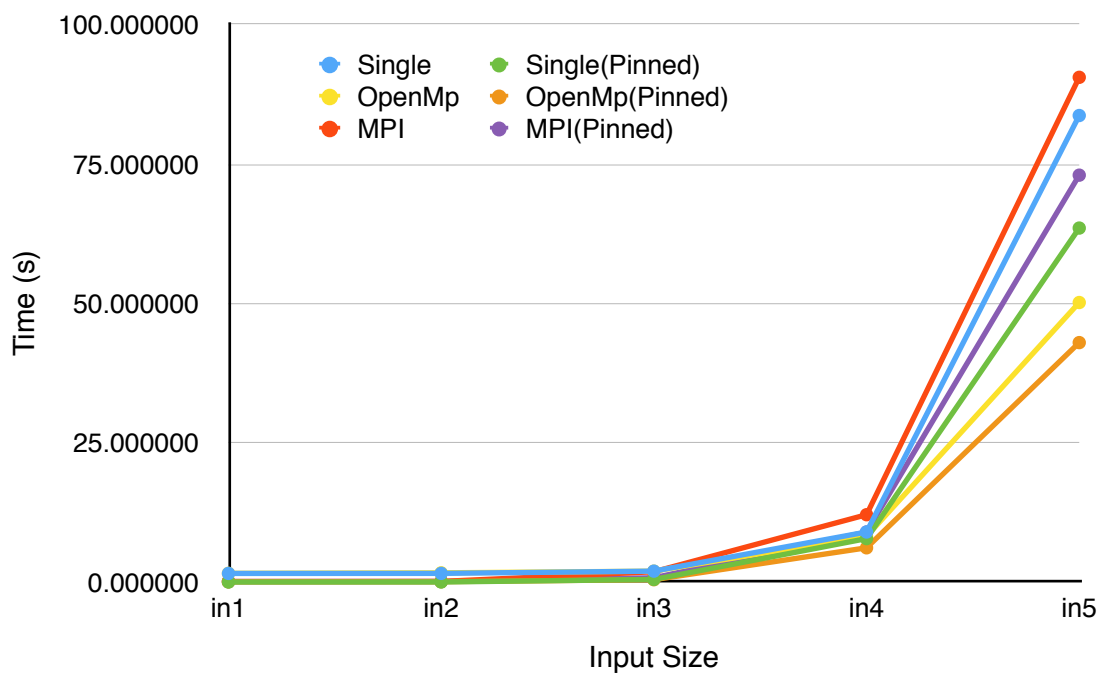
```

block(10, 10)
thread(10, 10)
B = 64
OpenMp (2 threads)
MPI (2 processes)

```

- Results :

	Single	Single (Pinned)	OpenMp	OpenMp (Pinned)	MPI	MPI (Pinned)
in1	1.545760	0.000492	1.562793	0.057982	0.096190	0.094115
in2	1.554278	0.008190	1.652941	0.072133	0.110442	0.110438
in3	1.957913	0.465849	1.953298	0.425449	1.846156	0.765513
in4	8.973730	7.785813	8.200664	6.138620	12.073054	8.966215
in5	83.645468	63.467112	47.118546	42.931893	90.494182	72.952075



- Observation :


- (1) Total execution time increases as the input size grows.
- (2) Pinned memory speeds up the execution to certain extent, but the effect of pinned memory is not apparent in perspective of total execution time.

- Optimization (Pinned Memory)

- Input Setting :

```
block(10, 10)
thread(10, 10)
B = 64
OpenMp (2 threads)
MPI (2 processes)
```

- Results :



- Observation :

- (1) From the perspective of memory, pinned memory improves the memory copy rate considerably. It copies the host memory variables to the pinned array so GPU can access the variable directly through the pinned memory.
- (2) For single version, pinned memory approximately triples the memory copy rate.
- (3) For OpenMp and MPI version, since they call the same amount of `cudaMemcpy()`, their results are similar.

Experience & Conclusion

Blocking method indeed provides a more efficient way to solve the APSP problem when size is large, but the relationship between number of block, thread, and the value of blocking factor decides the final performance. Blocking factor has its effective range as mentioned in the experiment, while the number of thread also contributes to the performance. Giving too few threads forces the threads to wait for others; giving too many threads makes some threads to be idling, which also affect the utilization of such resource.

In HW4, I learned how to do CUDA programming, and utilize the knowledge acquired from class. The concept of blocks and threads and all the indexing are actually really abstract to me. However, after the blocked APSP assignment, I became more familiar with the idea of CUDA programming. Tools like MPI and OpenMp also become more handy when doing this assignment after all these experience.

Pinned memory optimization puzzles me at first because no matter how many times I tried, the total execution time only decreases by very few seconds, and sometimes the execution even got longer. Then I realized comparison should be done in perspective of memory transfer, which shows more reasonable outcomes.

To me, blocked APSP is not only a brilliant, but also an algorithm that is hard to think of and implement. If the sequential code was not provided, I would still be doing the basic coding part of the assignment.